

CV IS Study

Donald Winkelman, Maxwell Tang, Jaime Davila

THE FUTURE!

1 Abstract

2 Background

Note: I am writing this assuming that the reader knows nothing about MAMBA or S4 and wants to learn everything that I learned

In this paper, we use a network architecture called MAMBA, which is centered around a type of sequence transformation called a state space model (SSM). State space models take some variable-length sequence of data (text, audio, even image data), and convert it into some other variable-length sequence of data, similarly to what a transformer does. As for how state space models calculate their output sequences, all state space models are discretizations of the following model, where x is the input sequence, and y is the output sequence:

$$\begin{aligned} n \in \mathbb{N} \quad A \in \mathbb{M}^{n \times n} \quad B \in \mathbb{M}^{n \times 1} \quad C \in \mathbb{M}^{1 \times n} \\ \frac{d\vec{s}}{dt} = Bx + A\vec{s} \\ y = C\vec{s} \end{aligned}$$

In English, state space models store an internal state that continuously changes w.r.t the input, and then compute the output at each timestep based on the internal state. Intuitively, \vec{s} is the internal state of the model. For a language model, \vec{s} might store the most recent sentence; For an image model, it might store nearby lines and shapes. Notice that this formula uses derivatives, meaning that it only applies to continuous sequences.

Since we can't store general continuous sequences in hardware, real implementations of SSMs use discretizations, with some finite timestep. This means that the actual dynamics of real SSMs are as follows:

$$\begin{aligned} \vec{s}_{t+1} &= \exp(\Delta A)\vec{s}_t + \Delta B\vec{x}_t \\ y_t &= C\vec{s}_{t+1} \end{aligned}$$

In addition, real models tend to work with multidimensional sequences, so each SSM layer is typically a stack of multiple SSMs.

2.1 hyperparameters

A common nomenclature for the hyperparameters of state space models is as follows:

- B - The batch size
- L - The length of the sequence
- D - The number of sequences per instance. This can be thought of as the number of "channels" that the layer takes as input.
- N - The state size for one SSM

To give an example, at any given time step, the total size of the state would be BDN , since we have B instances of SSMs, D separate states vectors within each SSM, N real number for each state vector.

2.2 S4

S4[1] is one implementation of state space models. S4 makes all parameters (A , B , C , and Δ) fixed within any evaluation step, and it enforces a special structure on A . Since each parameter is constant with respect to the position in the sequence, the overall transformation is invariant to translations. In addition, it can be shown that discrete SSMs are linear transformations of the input. This means that SSMs are equivalent to convolutions, and one of the insights of S4 is that this allows the model to be computed efficiently using fourier transforms. In addition, the special structure on A is that it is a normal plus low-rank (NPLR) matrix. This allows the authors to compute the convolutional kernel much more efficiently. Rather than requiring $\tilde{O}(LN^2)$ time to train as with direct computation, each S4 "channel" requires only $O(LN + N^2)$ time to train thanks to the

2.3 Mamba

3 Methodology

4 Results

5 Conclusion

References

- [1] Albert Gu, Karan Goel, and Christopher Ré. "Efficiently Modeling Long Sequences with Structured State Spaces". In: *The International Conference on Learning Representations (ICLR)*. 2022.